

# CS 446/ECE 449: Machine Learning

---

## Lecture 10: PAC Learning Theory (II)

Han Zhao  
02/15/2024



# Recap: Bayes Error

---

Bayes error rate:

$$\text{Bayes error: } \varepsilon_{\mu}^* := \inf_{f: \mathcal{X} \rightarrow \mathcal{Y}} \varepsilon_{\mu}(f)$$

Binary classification:

$$\text{Bayes error rate: } \varepsilon_{\mu}^* = \mathbb{E} \min \{ \Pr(Y = 1 | X), \Pr(Y = 0 | X) \}$$

$$\text{Bayes optimal classifier: } f_{\text{Bayes}}(X) := \begin{cases} 1 & \text{if } \Pr(Y = 1 | X) \geq \frac{1}{2} \\ 0 & \text{otherwise} \end{cases}$$

Regression with squared loss:

$$\text{Bayes error rate: } \varepsilon_{\mu}^* = \mathbb{E} \text{Var}[Y | X]$$

$$\text{Bayes optimal regressor: } f_{\text{Bayes}}(X) = \mathbb{E}[Y | X]$$

# Recap: Error Decomposition

Error decomposition:  $\forall f \in \mathcal{F}$ :

$$\varepsilon_{\mu}(f) = \underbrace{\varepsilon_{\mu}(f) - \inf_{f \in \mathcal{F}} \varepsilon_{\mu}(f)}_{\text{Estimation error}} + \underbrace{\inf_{f \in \mathcal{F}} \varepsilon_{\mu}(f) - \varepsilon_{\mu}^*}_{\text{Approximation error}} + \underbrace{\varepsilon_{\mu}^*}_{\text{Bayes error}}$$

**Estimation error**  
(depending on the size of our data and  $\mathcal{F}$ )

**Bayes error**  
(depending on the inherent noise in the data)

**Approximation error**  
(depending on the expressiveness of  $\mathcal{F}$ )

- Often the case, there is a trade-off between the estimation error and the approximation error
- If  $\mathcal{F}$  is more expressive, then the approximation error gets smaller but the estimation error gets larger
- If  $\mathcal{F}$  is more restricted, then the approximation error gets larger but the estimation error gets smaller (assume the size of training data is fixed)

# Lecture Today

---

- Probably Approximately Correct (PAC) framework
- Generalization analysis
- Vapnik–Chervonenkis dimension (VC dim)

# Probably Approximately Correct (PAC)

---

The learning process:

- We can choose a predictor  $f$  from some pre-defined class of functions  $\mathcal{F}$ , e.g., the class of linear predictors, decision trees, kernel machines, neural networks, etc.

We also have our training data  $\mathcal{D} := \{(x^{(i)}, y^{(i)})\}_{i=1}^n \sim \mu$  sampled independently and identically (iid) from the underlying distribution  $\mu$  over  $\mathcal{X} \times \mathcal{Y}$

We can then talk about two error measures (classification):

$$\text{Training error: } \hat{\varepsilon}_{\mathcal{D}}(f) := \frac{1}{n} \sum_{i=1}^n \mathbb{I}(f(x^{(i)}) \neq y^{(i)})$$

$$\text{Test error: } \varepsilon_{\mu}(f) := \mathbb{E}_{\mu} [\mathbb{I}(f(X) \neq Y)] = \Pr_{\mu}(f(X) \neq Y)$$

We are interested in finding  $f$  that minimizes the test error but we can only observe the training error

# Probably Approximately Correct (PAC)

---

For a given **hypothesis class**  $\mathcal{F}$ , can we relate the training and test errors?

Generalization error/gap:  $|\hat{\varepsilon}_{\mathcal{D}}(f) - \varepsilon_{\mu}(f)|$

Note:

- The generalization error is a random variable due to the randomness in  $\mathcal{D} \sim \mu$
- For any fixed  $f$ , we would expect the generalization error to be small:

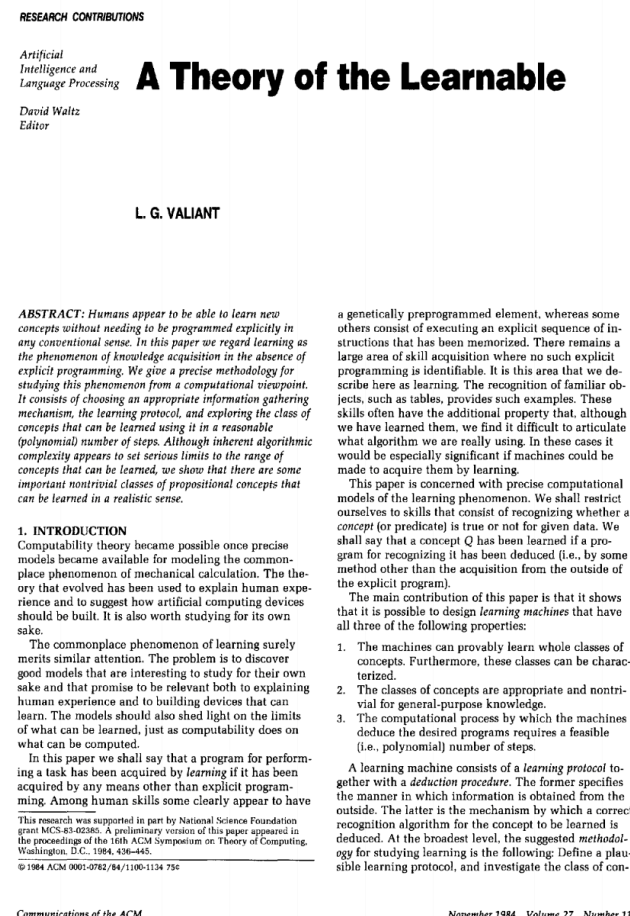
$$\mathbb{E} [\hat{\varepsilon}_{\mathcal{D}}(f)] = \mathbb{E} \left[ \frac{1}{n} \sum_{i=1}^n \mathbb{I}(f(x^{(i)}) \neq y^{(i)}) \right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E} [\mathbb{I}(f(x^{(i)}) \neq y^{(i)})] = \varepsilon_{\mu}(f)$$

The argument above is in expectation, and it does not necessarily apply to our specific training data  $\mathcal{D}$ . How about we consider a high-probability guarantee instead?

# Probably Approximately Correct (PAC)

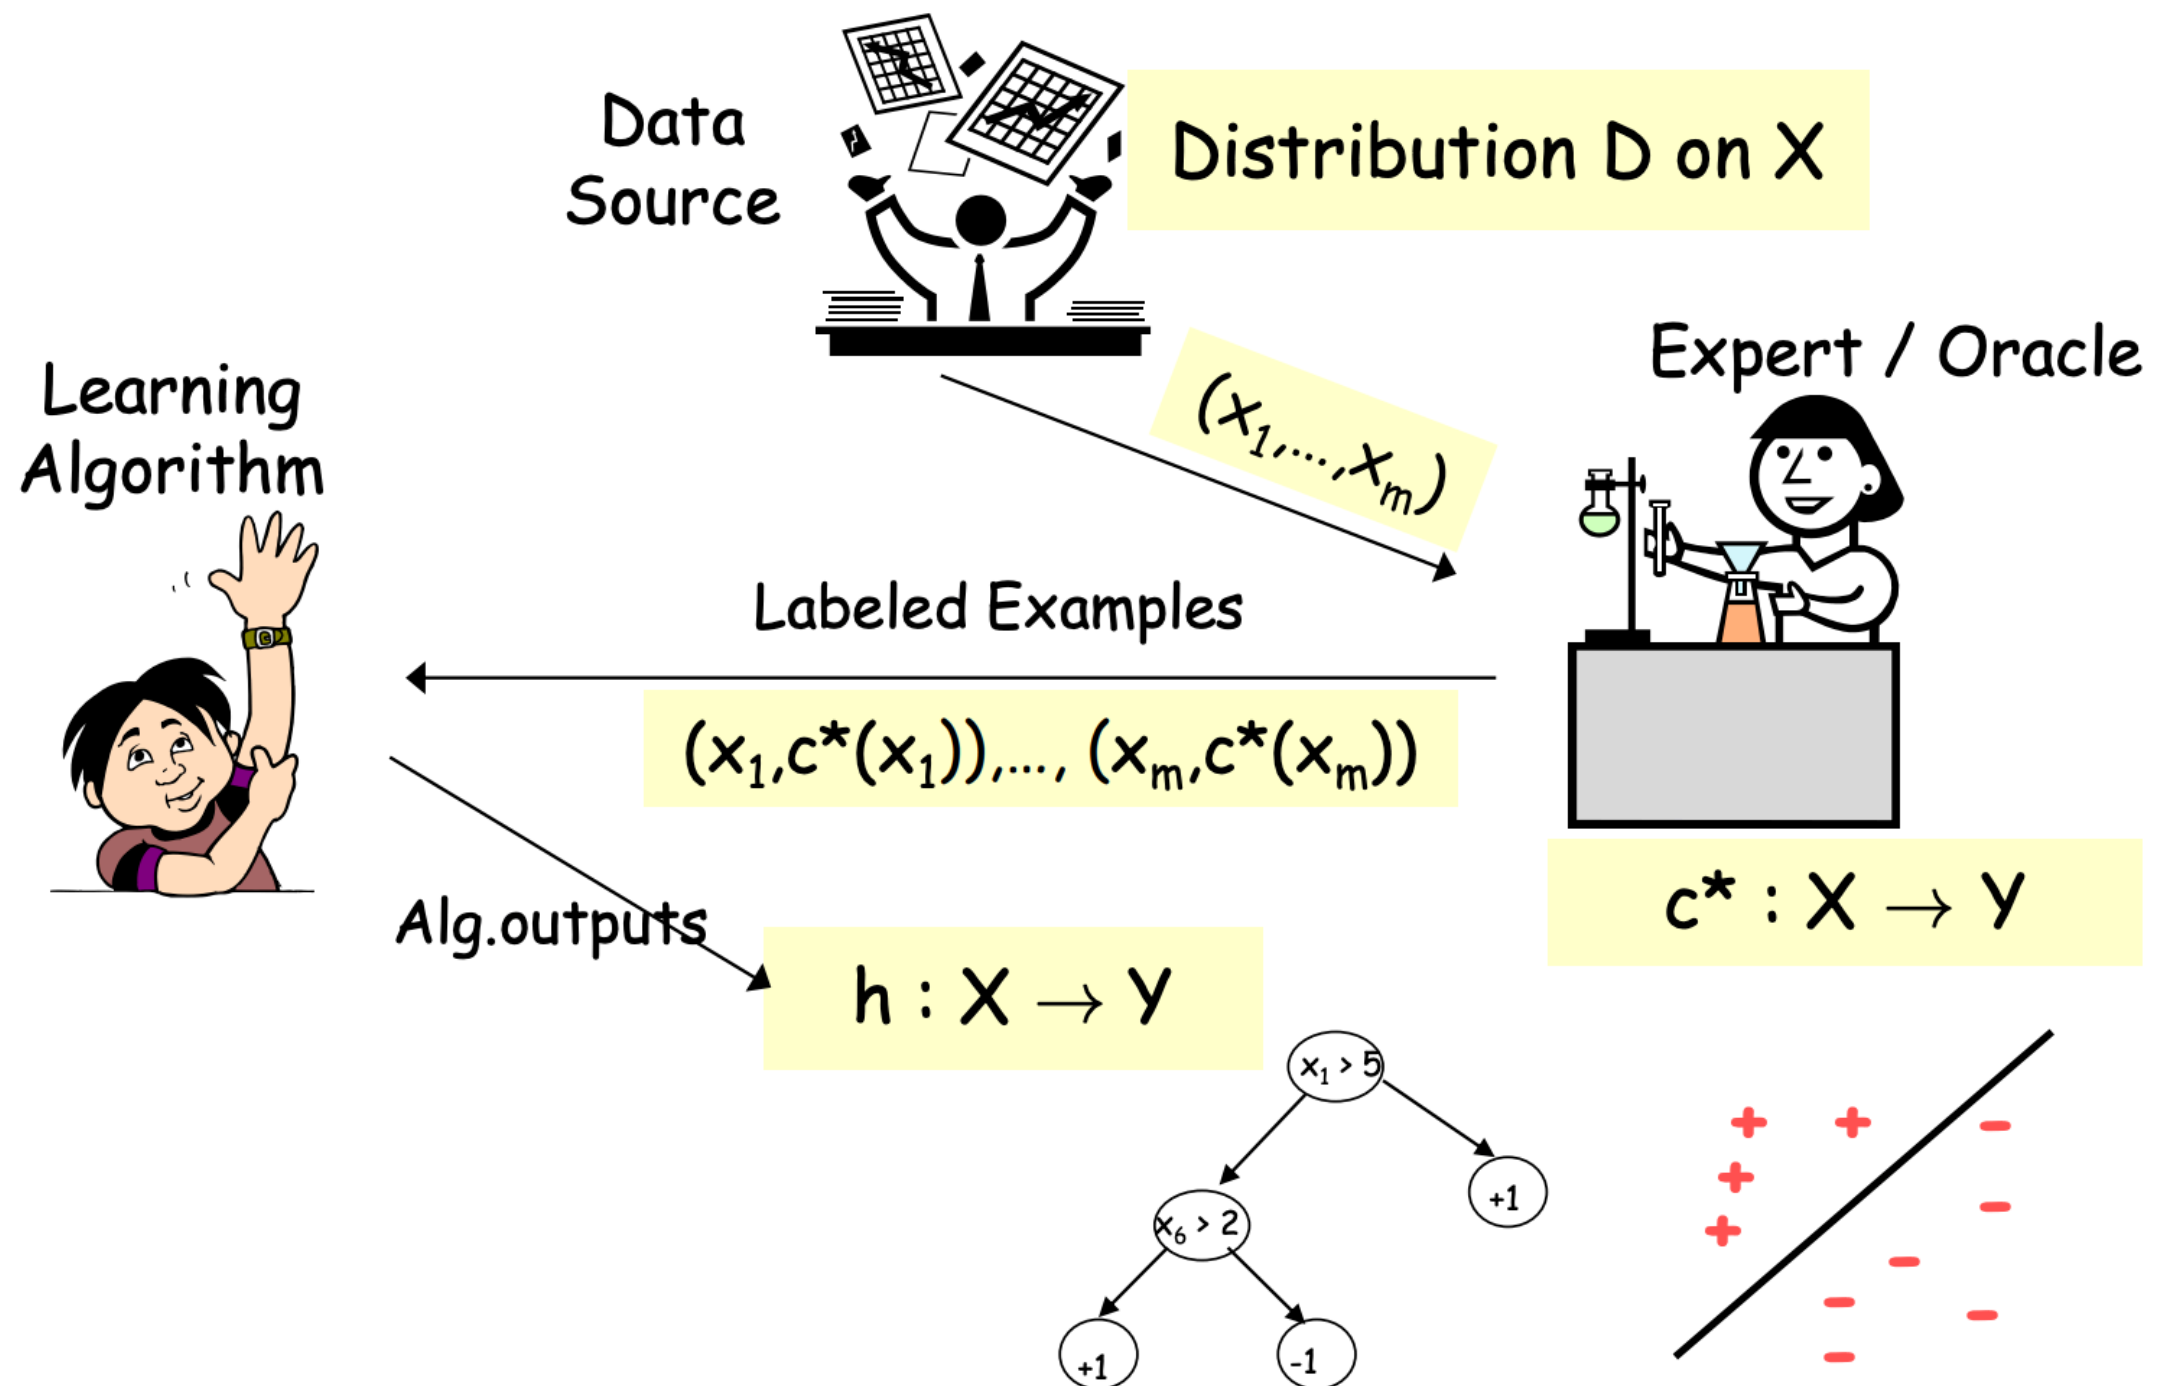
## Probably Approximately Correct (PAC, Valiant, CACM 1984)

- (Informal) A framework to quantify the meaning of learning a concept from samples
- With high probability (P), the learned predictor will have low generalization error (AC)
- No distributional assumption



# Probably Approximately Correct (PAC)

Probably Approximately Correct (PAC, Valiant, CACM 1984)





# Probably Approximately Correct (PAC)

---

Probably **A**pproximately **C**orrect (PAC, Valiant, CACM 1984)

Definition (PAC-learnable): A concept class  $\mathcal{F}$  is said to be PAC-learnable if there exists an algorithm  $\mathcal{A}$  such that for any  $0 < \epsilon, \delta < 1$ , for any distribution  $\mu$  over  $\mathcal{X}$  and for any target concept  $c \in \mathcal{F}$ , the following holds for any sample size  $n \geq \text{poly}(1/\epsilon, 1/\delta, d)$ :

$$\Pr_{\mathcal{D}}(\epsilon_{\mu}(f) \leq \epsilon) \geq 1 - \delta$$

where  $f$  is the output of the algorithm  $\mathcal{A}$ .

- $\epsilon$  is called the accuracy parameter
- $\delta$  is called the confidence parameter

# Probably Approximately Correct (PAC)

Probably Approximately Correct (PAC, Valiant, CACM 1984)

Definition (PAC-learnable): A concept class  $\mathcal{F}$  is said to be PAC-learnable if there exists an algorithm  $\mathcal{A}$  such that for any  $0 < \epsilon, \delta < 1$ , for any distribution  $\mu$  over  $\mathcal{X}$  and for any target concept  $c \in \mathcal{F}$ , the following holds for any sample size  $n \geq \text{poly}(1/\epsilon, 1/\delta, d)$ :


$$\Pr_{\mathcal{D}}(\epsilon_{\mu}(f) \leq \epsilon) \geq 1 - \delta$$

where  $f$  is the output of the algorithm  $\mathcal{A}$ .

- This holds for arbitrary target concept
- No assumption on the distribution  $\mu$
- PAC-learnability does not mention about the time complexity of running  $\mathcal{A}$
- The polynomial  $\text{poly}(1/\epsilon, 1/\delta, d)$  is called the sample complexity of  $\mathcal{A}$

# Generalization Analysis


A running example: learning with initial-segment



- $\mathcal{X} = (0,1)$
- $\mathcal{F} = \{c_a \in 2^{(0,1)} \mid c_a(x) = 1 \iff x \leq a\}$
- Let's consider a simple algorithm: return  $f = (\max_{x:x \in +} x + \min_{x:x \in -} x)/2$
- Assume the distribution over  $\mathcal{X}$  to be uniform

# Generalization Analysis

If the returned position is on the right of  $c$ :




- Error only happens at the interval between  $c$  and  $f$
- We want to upper bound the error probability:  $\Pr(\varepsilon_\mu(f) \geq \epsilon)$

# Generalization Analysis

---


If the returned position is on the right of  $c$ :



- Claim: there is no point in the training data from  $\mu$  that lies in this interval (?)

# Generalization Analysis

If the returned position is on the right of  $c$ :



$$\Pr(\varepsilon_\mu(f) \geq \epsilon \mid f \text{ on the right of } c)$$

$$\leq \Pr(\text{none of the training data lies in the interval})$$

$$\leq (1 - \epsilon)^n$$

$$\leq \exp(-n\epsilon)$$

iid assumption

$$\forall x, 1 - x \leq \exp(-x)$$

# Generalization Analysis

---

Similarly, if the returned position is on the left of  $c$ :

$$\Pr(\varepsilon_\mu(f) \geq \epsilon \mid f \text{ on the left of } c) \leq \exp(-n\epsilon)$$

Now, by a union bound ( $\Pr(A \cup B) \leq \Pr(A) + \Pr(B)$ ), and let  $L = f$  on the left of  $c$  and  $R = f$  on the right of  $c$

$$\begin{aligned}\Pr(\varepsilon_\mu(f) \geq \epsilon) &= \Pr(\varepsilon_\mu(f) \geq \epsilon \mid L) \Pr(L) + \Pr(\varepsilon_\mu(f) \geq \epsilon \mid R) \Pr(R) \\ &\leq \Pr(\varepsilon_\mu(f) \geq \epsilon \mid L) + \Pr(\varepsilon_\mu(f) \geq \epsilon \mid R) \\ &\leq 2 \exp(-n\epsilon) \\ &\leq \delta\end{aligned}$$

Solving for  $n$ , we get: it suffices if

$$n \geq \frac{1}{\epsilon} \log \frac{1}{\delta}$$

This shows that  $\mathcal{F}$  is PAC-learnable.

# Generalization Analysis

Could we generalize the previous results?

Realizable case with finite  $\mathcal{F}$ :  $|\mathcal{F}| < \infty, c \in \mathcal{F}$

**Theorem:** Let  $f$  be an empirical risk minimizer on a training data with  $n$  examples where

$$n \geq \frac{1}{\epsilon} \left( \log |\mathcal{F}| + \log \frac{1}{\delta} \right)$$

Then  $\Pr(\epsilon_\mu(f) \leq \epsilon) \geq 1 - \delta$ . Equivalently, with probability at least  $1 - \delta$ :

$$\epsilon_\mu(f) \leq \frac{1}{n} \left( \log |\mathcal{F}| + \log \frac{1}{\delta} \right)$$

Empirical Risk Minimization (ERM):

$$f_{\text{ERM}} = \mathcal{A}_{\text{ERM}}(\mathcal{D}) := \arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \mathbb{I}(f(x^{(i)}) \neq y^{(i)})$$

i.e., the ERM algorithm finds a predictor that minimizes the training loss



# Generalization Analysis

Realizable case with finite  $\mathcal{F}$ :  $|\mathcal{F}| < \infty, c \in \mathcal{F}$

**Theorem:** Let  $f$  be an empirical risk minimizer on a training data with  $n$  examples where

$$n \geq \frac{1}{\epsilon} \left( \log |\mathcal{F}| + \log \frac{1}{\delta} \right)$$

Then  $\Pr(\epsilon_\mu(f) \leq \epsilon) \geq 1 - \delta$ . Equivalently, with probability at least  $1 - \delta$ :

$$\epsilon_\mu(f) \leq \frac{1}{n} \left( \log |\mathcal{F}| + \log \frac{1}{\delta} \right)$$

- Fact: since we are using ERM under realizable case, the training error of the ERM solution will be 0
- Let's fix a classifier  $f$  and consider its true error. Instead, let  $f_{\text{ERM}}$  be the solution returned by the ERM algorithm

# Generalization Analysis

---

By definition of conditional probability:

$$\Pr \left( \hat{\varepsilon}_{\mathcal{D}}(f) = 0 \wedge \varepsilon_{\mu}(f) > \epsilon \right) \leq \Pr \left( \hat{\varepsilon}_{\mathcal{D}}(f) = 0 \mid \varepsilon_{\mu}(f) > \epsilon \right)$$

But,

$$\Pr \left( \hat{\varepsilon}_{\mathcal{D}}(f) = 0 \mid \varepsilon_{\mu}(f) > \epsilon \right) \leq (1 - \epsilon)^n$$

Hence, by union bound,

$$\Pr \left( \exists f \in \mathcal{F} : \hat{\varepsilon}_{\mathcal{D}}(f) = 0 \wedge \varepsilon_{\mu}(f) > \epsilon \right) \leq |\mathcal{F}| \cdot (1 - \epsilon)^n \leq \delta$$

On the other hand, we have

$$\begin{aligned} \Pr(\varepsilon_{\mu}(f_{\text{ERM}}) > \epsilon) &= \Pr(\hat{\varepsilon}_{\mathcal{D}}(f_{\text{ERM}}) = 0 \wedge \varepsilon_{\mu}(f_{\text{ERM}}) > \epsilon) \\ &\leq \Pr \left( \exists f \in \mathcal{F} : \hat{\varepsilon}_{\mathcal{D}}(f) = 0 \wedge \varepsilon_{\mu}(f) > \epsilon \right) \\ &\leq |\mathcal{F}| (1 - \epsilon)^n \leq |\mathcal{F}| \exp(-n\epsilon) \leq \delta \end{aligned}$$

Solving for  $n$ , we get

$$n \geq \frac{1}{\epsilon} \left( \log |\mathcal{F}| + \log \frac{1}{\delta} \right)$$

# Probably Approximately Correct (Agnostic)

So far we mainly talk about realizable case with finite hypothesis class.

Probably Approximately Correct (agnostic case)

Definition (PAC-learnable): A hypothesis space  $\mathcal{H}$  is said to be agnostic PAC-learnable if there exists an algorithm  $\mathcal{A}$  such that for any  $0 < \epsilon, \delta < 1$ , for all distribution  $\mu$  over  $\mathcal{X} \times \mathcal{Y}$ , the following holds for any sample size  $n \geq \text{poly}(1/\epsilon, 1/\delta, d)$ :

$$\Pr \left( \epsilon_{\mu}(f) \leq \min_{f' \in \mathcal{H}} \epsilon_{\mu}(f') + \epsilon \right) \geq 1 - \delta$$

where  $f$  is the output of the algorithm  $\mathcal{A}$ .

- $\epsilon$  is called the accuracy parameter
- $\delta$  is called the confidence parameter
- No assumption on  $\mu$  has been made
- Agnostic PAC-learnability does not mention about the time complexity of running  $\mathcal{A}$
- The polynomial  $\text{poly}(1/\epsilon, 1/\delta, d)$  is called the sample complexity of  $\mathcal{A}$

# Concentration Inequality

Some useful inequalities regarding the concentration of RVs

Theorem (Hoeffding's inequality): Let  $Z_1, \dots, Z_n$  be independent RVs where  $Z_i \in [a, b]$ . Then for any  $\epsilon > 0$ , the following inequality holds for the mean  $\bar{Z}_n = \frac{1}{n} \sum_{i=1}^n Z_i$ :

$$\Pr \left( \left| \bar{Z}_n - \mathbb{E}[\bar{Z}_n] \right| \geq \epsilon \right) \leq 2 \exp \left( -\frac{2n\epsilon^2}{(b-a)^2} \right)$$

Equivalent statement: with probability at least  $1 - \delta$ , we have:

$$\left| \bar{Z}_n - \mathbb{E}[\bar{Z}_n] \right| \leq (b-a) \sqrt{\frac{\log(2/\delta)}{2n}}$$

If  $Z_1, \dots, Z_n$  are iid, then  $\mathbb{E}[\bar{Z}_n] = \mathbb{E}[Z_i], \forall i \in [n]$  so we have

$$\left| \bar{Z}_n - \mathbb{E}[Z_1] \right| \leq (b-a) \sqrt{\frac{\log(2/\delta)}{2n}}$$

# Concentration Inequality

Some useful inequalities regarding the concentration of RVs

Theorem (Hoeffding's inequality): Let  $Z_1, \dots, Z_n$  be independent RVs where  $Z_i \in [a, b]$ . Then for any  $\epsilon > 0$ , the following inequality holds for the mean  $\bar{Z}_n = \frac{1}{n} \sum_{i=1}^n Z_i$ :

$$\Pr \left( \left| \bar{Z}_n - \mathbb{E}[\bar{Z}_n] \right| \geq \epsilon \right) \leq 2 \exp \left( -\frac{2n\epsilon^2}{(b-a)^2} \right)$$

Example: Coin flipping

- Suppose we have a coin with head probability  $p$
- We flipped the coin for 1000 times, with an average head frequency  $\hat{p}$
- How close will the frequency  $\hat{p}$  be to the true  $p$ ?

# Generalization Analysis

---

For any fixed  $f \in \mathcal{F}$ , we can use the Hoeffding's inequality to get a generalization bound:

Let  $Z_i = \mathbb{I}(f(X^{(i)}) \neq Y^{(i)}) \in \{0,1\} \subseteq [0,1]$ , then

**Training error:** 
$$\bar{Z}_n = \frac{1}{n} \sum_{i=1}^n Z_i = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(f(X^{(i)}) \neq Y^{(i)}) = \hat{\varepsilon}_{\mathcal{D}}(f)$$

**Test error:** 
$$\mathbb{E} [\bar{Z}_n] = \mathbb{E} [Z_1] = \varepsilon_{\mu}(f)$$

By Hoeffding's inequality, with probability at least  $1 - \delta$ , we have

$$\varepsilon_{\mu}(f) \leq \hat{\varepsilon}_{\mathcal{D}}(f) + \sqrt{\frac{\log(2/\delta)}{2n}}$$

**Note:** it is important to fix a predictor  $f$  in order for the analysis above to hold, i.e.,  $f$  cannot be the output of an algorithm  $\mathcal{A}$  that depends on the data  $\mathcal{D}$

# Generalization Analysis

---

What if  $f = \mathcal{A}(\mathcal{D})$ ?

Given  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n \sim \mu$  be a dataset of iid samples. Define our algorithm as follows:

$$f(x) := \begin{cases} y_i & \text{if } x = x_i \\ \text{"unknown"} & \text{otherwise} \end{cases}$$

Then  $\hat{\varepsilon}_{\mathcal{D}}(f) = 0$  and  $\varepsilon_{\mu}(f) = 1$ !

Why?

- Hoeffding's inequality cannot be applied anymore, since  $f$  is the outcome of an algorithm  $\mathcal{A}$  that depends on the data  $\mathcal{D}$ . In other words, given  $f$ , the data  $Z_i$  are no longer independent

Fix?

- Use a disjoint validation set to empirically estimate the error
- Pay a model complexity penalty term: with probability at least  $1 - \delta$ , for all  $f \in \mathcal{F}$  simultaneously, we have:

$$\varepsilon_{\mu}(f) \leq \hat{\varepsilon}_{\mathcal{D}}(f) + O\left(\sqrt{\frac{\text{complexity}(\mathcal{F}) + \log(1/\delta)}{n}}\right)$$

# Vapnik–Chervonenkis dimension (VC-dim)

---

## VC dimension:

Let  $\mathcal{F} : \mathbb{R}^d \rightarrow \{0,1\}$  be a set of binary functions. Then the VC dimension of  $\mathcal{F}$ , denoted by  $\text{VCdim}(\mathcal{F})$  is the cardinality of the largest set of points in  $\mathbb{R}^d$  that can be **shattered** by  $\mathcal{F}$ .

## Shattering:

Given a set  $\mathcal{D} \subseteq \mathbb{R}^d$  of size  $n$ , i.e.,  $|\mathcal{D}| = n$ , we say that  $\mathcal{D}$  can be shattered by  $\mathcal{F}$  iff

$$\forall S \subseteq \mathcal{D}, \exists f \in \mathcal{F} : \forall x \in S, f(x) = 1, \forall x \notin S, f(x) = 0$$

## Note:

- By definition, in order to claim the VC-dim of a given hypothesis class  $\mathcal{F}$  to be  $n$ , we need to verify the following two conditions:
  - ★  $\text{VCdim}(\mathcal{F}) \geq n$ :  $\exists \mathcal{D} \subseteq \mathbb{R}^d : |\mathcal{D}| = n, \mathcal{F}$  shatters  $\mathcal{D}$
  - ★  $\text{VCdim}(\mathcal{F}) \leq n$ :  $\forall \mathcal{D} \subseteq \mathbb{R}^d : |\mathcal{D}| = n + 1, \mathcal{D}$  cannot be shattered by  $\mathcal{F}$




# Vapnik–Chervonenkis dimension (VC-dim)


Example:  $d = 2, \mathcal{F} = \{\text{linear classifiers in } \mathbb{R}^2\}$

Claim:  $\text{VCdim}(\mathcal{F}) = 3$

Proof that  $\text{VCdim}(\mathcal{F}) \geq 3$ :



Proof that  $\text{VCdim}(\mathcal{F}) < 4$ : XOR



# Generalization Analysis

---

With VC dim as the complexity measure, we have the following **uniform** generalization bound:

Given  $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^n \sim \mu$  be a dataset of iid samples. Let  $\mathcal{F}$  be a hypothesis class of finite VC-dim, i.e.,  $\text{VCdim}(\mathcal{F}) < \infty$ , then for  $0 < \delta < 1$ , with probability at least  $1 - \delta$ , for **all**  $f \in \mathcal{F}$ :

$$\varepsilon_{\mu}(f) \leq \hat{\varepsilon}_{\mathcal{D}}(f) + O\left(\sqrt{\frac{\text{VCdim}(\mathcal{F}) + \log(1/\delta)}{n}}\right)$$

**Note:**

- As long as  $\text{VCdim}(\mathcal{F}) < \infty$ , as  $n \rightarrow \infty$ , we know that the training error converges to the test error
- The bound above gives the generalization error, and we can use the generalization error bound to provide an upper bound on the estimation error, i.e.,  $\varepsilon_{\mu}(f) - \inf_{f' \in \mathcal{F}} \varepsilon_{\mu}(f')$
- There are other forms of complexity measures to characterize the expressiveness/richness/powerfulness of a given hypothesis class, but it is beyond the scope of this course
- The bound above could be loose, i.e., the generalization error could be larger than 1 for classification problems

# Next Time

---

- Perceptron Algorithm
- Deep Learning